

1.	Prerequisites:.....	3
2.	Installation of Kafka – 90 Mins	5
3.	Installation Confluent Kafka (Local) – 30 Minutes.....	13
4.	Basic Kafka Operations - CLI (Topic) – 30 Mins	14
5.	Basic Kafka Operations - CLI (Producer & Consumer) – 30 Mins	20
6.	Kafka cluster – 90 Minutes	27
7.	Schema Registry - Manage Schemas for Topics – 30 Minutes	49
8.	Errors	56
I.	{test=LEADER_NOT_AVAILABLE} (org.apache.kafka.clients.NetworkClient)	56
9.	LOG verification + segment Sizing	58
10.	Annexure Code:	59
II.	DumplogSegment	59
III.	Settings	60
IV.	Resources	63

Hardware:

8 GB RAM , 30 GB HDD , Centos 7 or above OS. Access to internet.

Software Inventory:

- Zookeeper Version: apache-zookeeper-3.8.0-bin.tar
- Apache kafka : 2.13-3.2.1

2 Kafka – Administration

- JDK 11.0.16
- Eclipse for Linux. (Any Latest version for JEE Development)
- Status : Color is Verified

Last Updated: June 29 - 2023.

1. Prerequisites:

Using docker:

Steps to be performed:

- Instantiate a container, kafka0.
- You can copy files from the host to container using docker copy command.

Execute the following command, it will perform the following:

- Create a network : spark-net
- download the image, centos:8
- start a container with the name, kafka0 and mount host folder in /opt

```
#docker network create --driver bridge spark-net
```

```
#docker run --name kafka0 --hostname kafka0 -p 9092:9092 -p 8081:8081 -p 2181:2181 -p 9999:9999 -i -t --privileged --network spark-net -v /Users/henrypotsangbam/Documents/Software:/Software centos /usr/sbin/init
```

Note: **/Users/henrypotsangbam/Documents/Software** replace this directory with that of your system. Any folder to store software. E.x /tmp

You can verify the container from a separate terminal:

4 Kafka – Administration

```
(base) Henrys-MacBook-Air:~ henrypotsangbam$ docker ps
CONTAINER ID        IMAGE               COMMAND             CREATED            STATUS              PORTS
0945477d457b        ehenry0227/learning:spark-kafka-raw   "/bin/bash"        34 seconds ago    Up 33 seconds
ds                  0.0.0.0:6062->6062/tcp, 0.0.0.0:8081->8081/tcp
(base) Henrys-MacBook-Air:~ henrypotsangbam$
```

Perform the installation after this as a normal server.

Ends Here -----

Software:

- [kafka_2.13-3.2.1.tgz](#)
- jdk 11.0.1
- apache-zookeeper-3.6.3-bin.tar.gz
- confluent-7.2.2.tar

5 Kafka – Administration

2. Installation of Kafka – 90 Mins

Start kafka node and perform the following.

You need to install java before installing zookeeper and Kafka.

Installation of Java.

Extract jdk in /opt folder. Required software should be in /Software folder. Ensure to enter the appropriate jdk file of your system. Rename the jdk folder to jdk

```
#tar -xvf jdk-11.0.16_linux* -C /opt  
#cd /opt  
#mv jdk* jdk
```

Set the above path in the PATH variable and JAVA_HOME

Update profile as follow:

```
#vi ~/.bashrc  
  
export JAVA_HOME=/opt/jdk  
export PATH=$PATH:$JAVA_HOME/bin
```

Save the file and exit.

Update the shell scripts using the following command.

```
#bash
```

6 Kafka – Administration

Install Zookeeper

You can choose any of the options given below:

Option I (Fresh Installation): (We will use this for our lab)

The following steps install Zookeeper with a basic configuration in /opt/zookeeper.
Its configured to store data in /opt/data/zookeeper:

Extract the zookeeper archive file in /opt and rename the installation folder for brevity.

```
# tar -xvf apache-zookeeper* -C /opt  
# cd /opt  
#mv a*zookeeper* /opt/zookeeper  
#mkdir -p /opt/data/zookeeper
```

Create a zookeeper configuration file and update with the following values.

```
#vi /opt/zookeeper/conf/zoo.cfg  
tickTime=2000  
dataDir=/opt/data/zookeeper  
clientPort=2181
```

Update the zoo.cfg with the above entries. You can save the file using esc+wq!

7 Kafka – Administration

Start the zookeeper using the following scripts.

```
# /opt/zookeeper/bin/zkServer.sh start
```

```
[root@tos opt]# /opt/zookeeper/bin/zkServer.sh start
ZooKeeper JMX enabled by default
Using config: /opt/zookeeper/bin/.../conf/zoo.cfg
Starting zookeeper ... STARTED
[root@tos opt]#
```

Option II (Part of the Apache Kafka) – Skip.

```
#bin/zookeeper-server-start.sh config/zookeeper.properties
```

Option II Ends Here.

You can now validate that Zookeeper is running correctly in standalone mode by connecting to the client port and sending the four-letter command srvr:

```
#yum install telnet
```

```
#telnet localhost 2181
```

Enter the following in the zookeeper CLI.

```
srvr
```

8 Kafka – Administration

```
[root@tos opt]# telnet localhost 2181
Trying ::1...
Connected to localhost.
Escape character is '^].
srvr
Zookeeper version: 3.4.12-e5259e437540f349646870ea94dc2658c4e44b3b, built on 03/
27/2018 03:55 GMT
Latency min/avg/max: 0/0/0
Received: 2
Sent: 1
Connections: 1
Outstanding: 0
Zxid: 0x0
Mode: standalone
Node count: 4
Connection closed by foreign host.
[root@tos opt]#
```

After zookeeper installation, let us install the Kafka Broker.

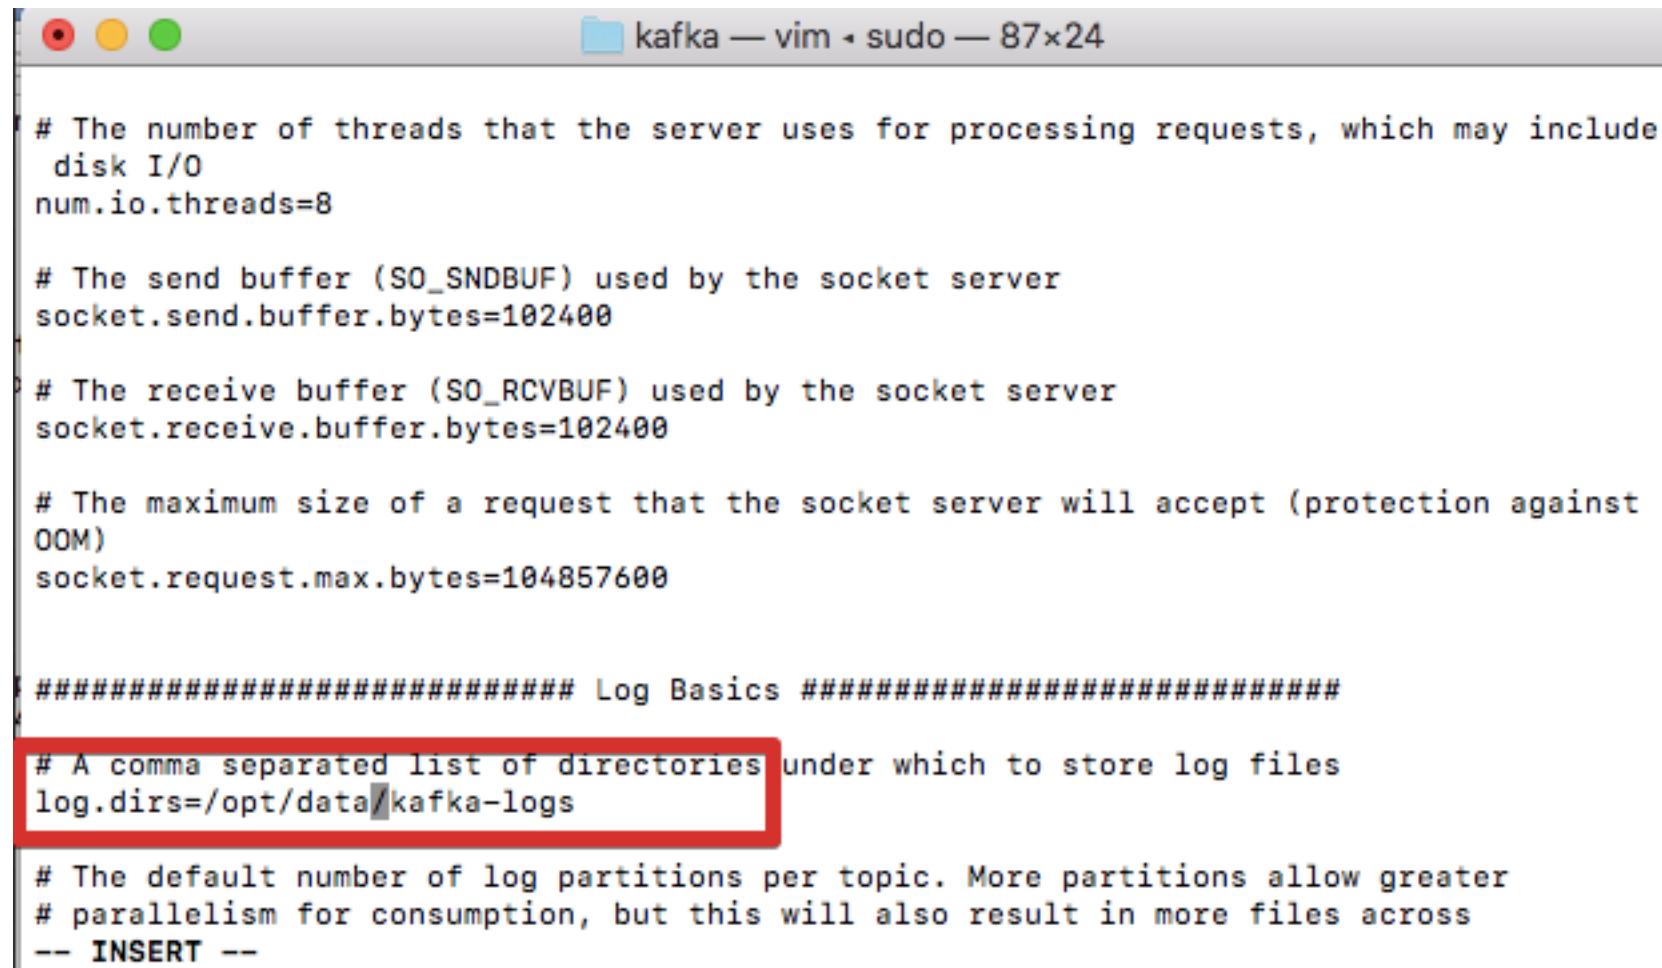
Installation of Kafka Broker

The following example installs Kafka in /opt/kafka, configured to use the Zookeeper server started previously and to store the message log segments stored in /tmp/kafka-logs: You need to execute the following command from the folder where the software is located.

```
# tar -xvf kafka* -C /opt
#cd /opt
# mv kafka* /opt/kafka
# mkdir /opt/data/kafka-logs
```

9 Kafka – Administration

Update `/opt/kafka/config/server.properties` to store Kafka Log in the above mention folder.



```
# The number of threads that the server uses for processing requests, which may include
# disk I/O
num.io.threads=8

# The send buffer (SO_SNDBUF) used by the socket server
socket.send.buffer.bytes=102400

# The receive buffer (SO_RCVBUF) used by the socket server
socket.receive.buffer.bytes=102400

# The maximum size of a request that the socket server will accept (protection against
# OOM)
socket.request.max.bytes=104857600

#####
##### Log Basics #####
#####

# A comma separated list of directories under which to store log files
log.dirs=/opt/data/kafka-logs

# The default number of log partitions per topic. More partitions allow greater
# parallelism for consumption, but this will also result in more files across
-- INSERT --
```

Back Up the configuration for later.

```
#cp /opt/kafka/config/server.properties /opt/kafka/config/server.properties_plain
```

Start the broker with the following command

```
# /opt/kafka/bin/kafka-server-start.sh -daemon /opt/kafka/config/server.properties
```

```
[root@tos opt]# /opt/kafka/bin/kafka-server-start.sh -daemon /opt/kafka/config/server.properties
[root@tos opt]# jps
3476 Kafka
3499 Jps
2895 QuorumPeerMain
[root@tos opt]#
```

```
#mkdir /opt/scripts
```

All the common execution scripts will be stored in the above folder.

The following scripts will start a zookeeper along with a broker. Create the following files and update with the following scripts. It will start the zookeeper and kafka broker using the mention script.

```
#cd /opt/scripts
#vi startABroker.sh
```

Enter the following in the above file.

```
##### Scripts Begin #####
#!/usr/bin/env bash
# Start Zookeeper.
```

```
/opt/zookeeper/bin/zkServer.sh start  
#Start Kafka Server  
/opt/kafka/bin/kafka-server-start.sh -daemon /opt/kafka/config/server.properties  
echo "Started Successfully"  
##### Scripts End #####
```

To shutdown the Broker, find the process and kill.

```
ps -eaf | grep java
```

or

create a script:

```
#vi /opt/scripts/stopBroker.sh
```

with the following commands in it.

```
#!/usr/bin/env bash  
/opt/kafka/bin/kafka-server-stop.sh
```

To Stop Zookeeper create the following script. Don't include the ---- line.

```
#vi /opt/scripts/stopZookeeper.sh
```

Update the following commands in the above script and save it.

```
#!/usr/bin/env bash  
# Stop Zookeeper.  
/opt/zookeeper/bin/zkServer.sh stop
```

```
echo "Stop zookeeper Successfully"
```

To reinitialize the Cluster.

```
#vi reinitializeCluster.sh
```

```
#!/usr/bin/env bash
rm -fr /opt/kafka/data/kafka-logs/*
rm -fr /opt/kafka/data/zookeeper/*
cp /opt/kafka/config/server.properties_plain /opt/kafka/config/server.properties
echo "Reinitialize Successfully"
```

Lab Installation completes End here.

3. Installation Confluent Kafka (Local) – 30 Minutes

The purpose of this lab is to demonstrate the basic and most powerful capabilities of Confluent Platform – Schema Registry.

Install Confluent kafka with the following step.

Inflate the confluent kafka compress file as shown below. Software should be in /Software

```
#tar -xvf confluent-7.2.2.tar -C /opt
```

Rename the folder.

```
#cd /opt  
#mv confluent-7.2.2 confluent
```

Set the environment variable for the Confluent Platform directory.
`export CONFLUENT_HOME=/opt/confluent`

Set your PATH variable:

```
# vi ~/.bashrc  
    export PATH=/opt/confluent/bin:${PATH};
```

type the following to activate the script.

```
#bash
```

----- Lab Ends Here -----

4. Basic Kafka Operations - CLI (Topic) – 30 Mins

In this lab you will be able to create a topic and perform some operations to understand the information about topic like partition and replication.

You need to start the broker using startABroker.sh. The script should be in /opt/scripts folder

```
#sh startABroker.sh
```

```
#jps
```

```
[root@tos scripts]# sh startABroker.sh
ZooKeeper JMX enabled by default
Using config: /opt/zookeeper/bin/../conf/zoo.cfg
Starting zookeeper ... STARTED
Started Successfully
[root@tos scripts]# jps
11665 Jps
11646 Kafka
11375 QuorumPeerMain
[root@tos scripts]#
```

Once the Kafka broker is started, we can verify that it is working by performing some simple operations against the broker; creating a test topic etc.

Create and verify details about topic:

```
# /opt/kafka/bin/kafka-topics.sh --create --bootstrap-server kafka0:9092 --
replication-factor 1 --partitions 1 --topic test
```

```
[root@kafka0 scripts]# /opt/kafka/bin/kafka-topics.sh --create --bootstrap-server kafka0:9092 --replication-factor 1 --partitions 1 --topic test
Created topic test.
```

```
# /opt/kafka/bin/kafka-topics.sh --create --bootstrap-server kafka0:9092 --replication-factor 1 --partitions 12 --topic IBM

# /opt/kafka/bin/kafka-topics.sh --bootstrap-server kafka0:9092 --describe --topic test

#/opt/kafka/bin/kafka-topics.sh --bootstrap-server kafka0:9092 --describe --topic IBM
```

```
[root@kafka0 scripts]# /opt/kafka/bin/kafka-topics.sh --create --bootstrap-server kafka0:9092 --replication-factor 1 --partitions 12 --topic IBM
Created topic IBM.
[root@kafka0 scripts]# /opt/kafka/bin/kafka-topics.sh --bootstrap-server kafka0:9092 --describe --topic test
Topic: test    TopicId: 9Rreu1aUR6uHfz3GhNYhtA PartitionCount: 1      ReplicationFactor: 1      Configs: segment.bytes=1073741824
  Topic: test    Partition: 0    Leader: 0      Replicas: 0      Isr: 0
[root@kafka0 scripts]# /opt/kafka/bin/kafka-topics.sh --bootstrap-server kafka0:9092 --describe --topic IBM
Topic: IBM    TopicId: eY3Tl3N_T6iYY_ZzV70Pjg PartitionCount: 12     ReplicationFactor: 1      Configs: segment.bytes=1073741824
  Topic: IBM    Partition: 0    Leader: 0      Replicas: 0      Isr: 0
  Topic: IBM    Partition: 1    Leader: 0      Replicas: 0      Isr: 0
  Topic: IBM    Partition: 2    Leader: 0      Replicas: 0      Isr: 0
  Topic: IBM    Partition: 3    Leader: 0      Replicas: 0      Isr: 0
  Topic: IBM    Partition: 4    Leader: 0      Replicas: 0      Isr: 0
  Topic: IBM    Partition: 5    Leader: 0      Replicas: 0      Isr: 0
  Topic: IBM    Partition: 6    Leader: 0      Replicas: 0      Isr: 0
  Topic: IBM    Partition: 7    Leader: 0      Replicas: 0      Isr: 0
  Topic: IBM    Partition: 8    Leader: 0      Replicas: 0      Isr: 0
  Topic: IBM    Partition: 9    Leader: 0      Replicas: 0      Isr: 0
  Topic: IBM    Partition: 10   Leader: 0      Replicas: 0      Isr: 0
  Topic: IBM    Partition: 11   Leader: 0      Replicas: 0      Isr: 0
[root@kafka0 scripts]#
```

Verify the no of partition in the output above.

list and describe topic.

```
#/opt/kafka/bin/kafka-topics.sh --bootstrap-server kafka0:9092 --list
```

```
[root@kafka0 scripts]# /opt/kafka/bin/kafka-topics.sh --bootstrap-server kafka0:9092 --list
IBM
test
[root@kafka0 scripts]#
```

List and describe Topics

What does the tool do?

This tool lists the information for a given list of topics. If no topics are provided in the command line, the tool queries zookeeper to get all the topics and lists the information for them. The fields that the tool displays are - topic name, partition, leader, replicas, isr.

How to use the tool?

List only single topic named "test" (prints only topic name)

```
#/opt/kafka/bin/kafka-topics.sh --list --bootstrap-server kafka0:9092 --topic test
```

List all topics (prints only topic names)

```
#/opt/kafka/bin/kafka-topics.sh --list --bootstrap-server kafka0:9092
```

```
[root@kafka0 scripts]# /opt/kafka/bin/kafka-topics.sh --list --bootstrap-server kafka0:9092 --topic test
test
[root@kafka0 scripts]# /opt/kafka/bin/kafka-topics.sh --list --bootstrap-server kafka0:9092
IBM
test
[root@kafka0 scripts]# █
```

Describe only single topic named "test" (prints details about the topic)

```
#/opt/kafka/bin/kafka-topics.sh --describe --bootstrap-server kafka0:9092 --topic
test
```

Describe all topics (prints details about the topics)

```
#/opt/kafka/bin/kafka-topics.sh --describe --bootstrap-server kafka0:9092
```

```
[root@kafka0 scripts]# /opt/kafka/bin/kafka-topics.sh --describe --bootstrap-server kafka0:9092 --topic test
Topic: test      TopicId: 9Rreu1aUR6uHfz3GHnYhtA PartitionCount: 1      ReplicationFactor: 1      Configs: segment.bytes=1073741824
    Topic: test      Partition: 0      Leader: 0      Replicas: 0      Isr: 0
[root@kafka0 scripts]# /opt/kafka/bin/kafka-topics.sh --describe --bootstrap-server kafka0:9092
Topic: test      TopicId: 9Rreu1aUR6uHfz3GHnYhtA PartitionCount: 1      ReplicationFactor: 1      Configs: segment.bytes=1073741824
    Topic: test      Partition: 0      Leader: 0      Replicas: 0      Isr: 0
Topic: IBM      TopicId: eY3Tl3N_T6iYY_ZzV70Pjg PartitionCount: 12      ReplicationFactor: 1      Configs: segment.bytes=1073741824
    Topic: IBM      Partition: 0      Leader: 0      Replicas: 0      Isr: 0
    Topic: IBM      Partition: 1      Leader: 0      Replicas: 0      Isr: 0
    Topic: IBM      Partition: 2      Leader: 0      Replicas: 0      Isr: 0
    Topic: IBM      Partition: 3      Leader: 0      Replicas: 0      Isr: 0
    Topic: IBM      Partition: 4      Leader: 0      Replicas: 0      Isr: 0
    Topic: IBM      Partition: 5      Leader: 0      Replicas: 0      Isr: 0
    Topic: IBM      Partition: 6      Leader: 0      Replicas: 0      Isr: 0
    Topic: IBM      Partition: 7      Leader: 0      Replicas: 0      Isr: 0
    Topic: IBM      Partition: 8      Leader: 0      Replicas: 0      Isr: 0
    Topic: IBM      Partition: 9      Leader: 0      Replicas: 0      Isr: 0
    Topic: IBM      Partition: 10     Leader: 0      Replicas: 0      Isr: 0
    Topic: IBM      Partition: 11     Leader: 0      Replicas: 0      Isr: 0
[root@kafka0 scripts]#
```

We will understand the output in details later.

Create Topics

What does the tool do?

By default, Kafka auto creates topic if "auto.create.topics.enable" is set to true on the server. This creates a topic with a default number of partitions, replication factor and uses Kafka's default scheme to do replica assignment. Sometimes, it may be required that we would like to customize a topic while creating it. This tool helps to create a topic and also specify the number of partitions, replication factor and replica assignment list for the topic.

How to use the tool?

create topic with default settings

```
#/opt/kafka/bin/kafka-topics.sh --create --bootstrap-server kafka0:9092 --topic topic1 --partitions 2 --replication-factor 1
```

```
[root@tos scripts]# /opt/kafka/bin/kafka-topics.sh --create --zookeeper localhost:2181 --topic topic1 --partitions 2 --replication-factor 1
Created topic "topic1".
[root@tos scripts]#
```

Create a topic with replication 2.

```
#/opt/kafka/bin/kafka-topics.sh --create --bootstrap-server kafka0:9092 --topic topic2 --partitions 2 --replication-factor 2
```

```
[root@kafka0 scripts]# /opt/kafka/bin/kafka-topics.sh --create --bootstrap-server kafka0:9092 --topic topic1 --partitions 2 --replication-factor 1
Created topic topic1.
[root@kafka0 scripts]# /opt/kafka/bin/kafka-topics.sh --create --bootstrap-server kafka0:9092 --topic topic2 --partitions 2 --replication-factor 2
Error while executing topic command : Replication factor: 2 larger than available brokers: 1.
[2023-07-03 15:53:18,020] ERROR org.apache.kafka.common.errors.InvalidReplicationFactorException: Replication factor: 2 larger than available brokers: 1.
(kafka.admin.TopicCommand$)
[root@kafka0 scripts]#
```

As shown above, it generates an error. Since there is only a single broker. It will fix later.

-----Lab CLI completes End here-----

5. Basic Kafka Operations - CLI (Producer & Consumer) – 30 Mins

In this lab we will send message to the broker and consumer message using the kafka inbuilt commands.

You need to complete the previous lab before proceeding ahead.

You need to start the broker using startABroker.sh if not done earlier. The script should be in /opt/scripts folder

```
#sh startABroker.sh  
#jps
```

```
[root@tos scripts]# sh startABroker.sh  
ZooKeeper JMX enabled by default  
Using config: /opt/zookeeper/bin/../conf/zoo.cfg  
Starting zookeeper ... STARTED  
Started Successfully  
[root@tos scripts]# jps  
11665 Jps  
11646 Kafka  
11375 QuorumPeerMain  
[root@tos scripts]#
```

Sent message to **test** topic: Open a console to send message to the topic, test. Enter some text as shown below.

```
# /opt/kafka/bin/kafka-console-producer.sh --broker-list kafka0:9092 --topic test
```

```
Test Message 1
```

```
Test Message 2
```

```
^D
```

```
#
```

```
[root@kafka0 scripts]# [root@kafka0 scripts]# /opt/kafka/bin/kafka-console-producer.sh --broker-list kafka0:9092 --topic test
>Test Message 1
>Test Message 2
>
```

Consume messages from a test topic: As soon as you enter the following script in a separate terminal, you should be able to consume the messages that we have type in the producer console.

```
# /opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server kafka0:9092 --topic
test --from-beginning
```

```
[root@kafka0 scripts]# [root@kafka0 scripts]# /opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server kafka0:9092 --topic test --from-beginning
Test Message 1
Test Message 2
>
```

DumpLogSegment

You can also view the message publish to a topic from a log file.

```
/opt/kafka/bin/kafka-run-class.sh kafka.tools.DumpLogSegments --deep-iteration --
print-data-log --files \
/opt/data/kafka-logs/test-0/000000000000000000.log | head -n 12
```

```
[root@kafka0 scripts]# /opt/kafka/bin/kafka-run-class.sh kafka.tools.DumpLogSegments --deep-iteration --print-data-log --files /opt/data/kafka-logs/test-0/00
0000000000000000.log | head -n 12
Dumping /opt/data/kafka-logs/test-0/000000000000000000.log
Starting offset: 0
baseOffset: 0 lastOffset: 0 count: 1 baseSequence: 0 lastSequence: 0 producerId: 0 producerEpoch: 0 partitionLeaderEpoch: 0 isTransactional: false isControl:
false deleteHorizonMs: OptionalLong.empty position: 0 CreateTime: 1688379972311 size: 82 magic: 2 compresscodec: none crc: 1391971847 isinvalid: true
| offset: 0 CreateTime: 1688379972311 keySize: -1 valueSize: 14 sequence: 0 headerKeys: [] payload: Test Message 1
baseOffset: 1 lastOffset: 1 count: 1 baseSequence: 1 lastSequence: 1 producerId: 0 producerEpoch: 0 partitionLeaderEpoch: 0 isTransactional: false isControl:
false deleteHorizonMs: OptionalLong.empty position: 82 CreateTime: 1688379975484 size: 82 magic: 2 compresscodec: none crc: 2251353974 isinvalid: true
| offset: 1 CreateTime: 1688379975484 keySize: -1 valueSize: 14 sequence: 1 headerKeys: [] payload: Test Message 2
[root@kafka0 scripts]#
```

Determine the offset of the messages.

CLI to list offset number.

```
#/opt/kafka/bin/kafka-run-class.sh kafka.tools.GetOffsetShell --broker-list
kafka0:9092 --topic test
```

```
[root@kafka0 scripts]# /opt/kafka/bin/kafka-run-class.sh kafka.tools.GetOffsetShell --broker-list kafka0:9092 --topic test
test:0:2
[root@kafka0 scripts]#
```

Partition 0 offset 2.

Sending Message with Key and Value

```
#/opt/kafka/bin/kafka-console-producer.sh --broker-list kafka0:9092 --topic IBM --property "parse.key=true" --property "key.separator=:"
```

Enter
WAS:600
DB2:700
ESB:900

```
[root@kafka0 scripts]# /opt/kafka/bin/kafka-console-producer.sh --broker-list kafka0:9092 --topic IBM --property "parse.key=true" --property "key.separator=:"  
>WAS:600  
>DB2:700  
>ESB:900  
>|
```

```
#/opt/kafka/bin/kafka-console-consumer.sh --topic IBM --bootstrap-server kafka0:9092 --from-beginning \  
--property print.key=true \  
--property key.separator=":"
```

```
[root@kafka0 /]# /opt/kafka/bin/kafka-console-consumer.sh --topic IBM --bootstrap-server kafka0:9092 --from-beginning \  
> --property print.key=true \  
> --property key.separator=":"  
WAS:600  
ESB:900  
DB2:700  
|
```

Let us understand consumer group load balancing concepts.

- Create a topic with 2 partitions – **city_info**.
- Initiate 2 consumers with same consumer Group
- Send 6 messages to the topic
- Verify load get distributed among the 2 consumers.

```
#/opt/kafka/bin/kafka-topics.sh --create --bootstrap-server kafka0:9092 --topic city_info --partitions 2 --replication-factor 1
```

Initiate 2 consumers with same consumer Group

Open in a separate 2 terminals.

```
#/opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server kafka0:9092 --topic city_info --consumer-property group.id=citygroup
```

```
#/opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server kafka0:9092 --topic city_info --consumer-property group.id=citygroup
```

- Produce about 6 messages to the topic

```
#/opt/kafka/bin/kafka-console-producer.sh --broker-list kafka0:9092 --topic city_info
```

```
[root@kafka0 scripts]# /opt/kafka/bin/kafka-console-producer.sh --broker-list kafka0:9092 --topic city_info
>Imphal is scenic
>Delhi is classic
>Mumbai is cosmopolitan
>Bangalore is cool
>Dubai is warm
>Bangkok is wow
>
```

- load get distributed among the 2 consumers.

```
@kafka0:/ (com.docker.cli) Last login: Tue Jul 4 08:42:24 on ttys000 henrypotsangbam@Henrys-MacBook-Pro ~ % docker exec -it kafka0 bash [root@kafka0 /]# /opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server kafka0:9092 --topic city_info --consumer-property group.id=citygroup
Delhi is classic
Bangalore is cool
Bangkok is wow
```

```
@kafka0:/ (com.docker.cli) Last login: Tue Jul 4 08:54:34 on ttys001 henrypotsangbam@Henrys-MacBook-Pro ~ % docker exec -it kafka0 bash [root@kafka0 /]# clear
[root@kafka0 /]#
[root@kafka0 /]# /opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server kafka0:9092 --topic city_info --consumer-property group.id=citygroup
Imphal is scenic
Mumbai is cosmopolitan
Dubai is warm
```

-----Lab CLI ends here-----

6. Kafka cluster – 90 Minutes

In this tutorial, we are going to run many Kafka Nodes on our system to Understand Kafka Failover in a cluster environment. So, you will need at least 8 GB of RAM in your system. You can run just two servers if you have less memory than 8 GB.

We are going to create a replicated topic and then demonstrate consumer group coordination along with broker failover. We will also demonstrate load balancing of Kafka consumers. We show how, with many groups, Kafka acts like a Publish/Subscribe. But, when we put all of our consumers in the same group, Kafka will load share the messages to the consumers in the same group (more like a queue than a topic in a traditional MOM sense).

If not already running, start [ZooKeeper](#).

Also, shut down Kafka from the first tutorial.

Next, you need to copy server properties for three brokers (detailed instructions to follow). Then we will modify these Kafka server properties to add unique Kafka ports, Kafka log locations, and unique Broker ids.

Then we will create three scripts to start these servers up using these properties, and then start the servers.

Lastly, we create replicated topic and use it to demonstrate Kafka consumer failover, and Kafka broker failover.

Create three new Kafka server-n.properties files

In this section, we will perform the following steps:

- copy the existing Kafka `server.properties` to `server-0.properties`, `server-1.properties`, and `server-2.properties`.
- Update `server-0.properties` to set `log.dirs` to “`/opt/data/kafka-logs/kafka-0`”.
- Modify `server-1.properties` to set `port` to `9093`, broker `id` to `1`, and `log.dirs` to “`/opt/data/kafka-logs/kafka-1`”.
- Lastly modify `server-2.properties` to use port `9094`, broker `id` `2`, and `log.dirs` “`/opt/data/kafka-logs/kafka-2`”.

Copy server properties file as follows: We will store all server's configuration in a single folder config - `kafka-config/config`.

```
$ cd /opt
$ mkdir -p kafka-config/config
$ cp kafka/config/server.properties kafka-config/config/server-0.properties
$ cp kafka/config/server.properties kafka-config/config/server-1.properties
$ cp kafka/config/server.properties kafka-config/config/server-2.properties
```

```
[root@tos opt]# mkdir -p kafka-config/config
[root@tos opt]# pwd
/opt
[root@tos opt]# ls
couchbase  data  jdk1.8.0_45  kafka  kafka-config  zookeeper  zookeeper.out
[root@tos opt]# cp kafka/config/server.properties kafka-config/server-0.properties
[root@tos opt]# ls
couchbase  data  jdk1.8.0_45  kafka  kafka-config  zookeeper  zookeeper.out
[root@tos opt]# cp kafka/config/server.properties kafka-config/config/server-1.properties
[root@tos opt]# cp kafka/config/server.properties kafka-config/config/server-2.properties
[root@tos opt]#
```

With your favourite text editor update server-0.properties so that `log.dirs` is set to `./logs/kafka-0`. Leave the rest of the file as it is. Make sure `log.dirs` is only defined once.

```
#vi /opt/kafka-config/config/server-0.properties
broker.id=0
listeners=PLAINTEXT://kafka0:9092
advertised.listeners=PLAINTEXT://kafka0:9092
log.dirs=/opt/data/kafka-logs/kafka-0
```

With your favorite text editor change `log.dirs`, `broker.id` and `log.dirs` of `server-1.properties` as follows.

```
#vi /opt/kafka-config/config/server-1.properties
broker.id=1
listeners=PLAINTEXT://kafka0:9093
advertised.listeners=PLAINTEXT://kafka0:9093
log.dirs=/opt/data/kafka-logs/kafka-1
```

With your favorite text editor change `log.dirs`, `broker.id` and `log.dirs` of `server-2.properties` as follows.

```
#vi /opt/kafka-config/config/server-2.properties
broker.id=2
listeners=PLAINTEXT://kafka0:9094
advertised.listeners=PLAINTEXT://kafka0:9094
log.dirs=/opt/data/kafka-logs/kafka-2
```

Create Start-up scripts for these three Kafka servers

The startup scripts will just run `kafka-server-start.sh` with the corresponding properties file.

```
#vi /opt/scripts/startB0.sh
```

Copy the following in the above file.

```
#!/usr/bin/env bash
## Run Kafka
/opt/kafka/bin/kafka-server-start.sh "/opt/kafka-config/config/server-
0.properties"
```

```
#vi /opt/scripts/startB1.sh
#!/usr/bin/env bash
## Run Kafka
/opt/kafka/bin/kafka-server-start.sh "/opt/kafka-config/config/server-
1.properties"
```

```
#vi /opt/scripts/startB2.sh
#!/usr/bin/env bash
## Run Kafka
/opt/kafka/bin/kafka-server-start.sh "/opt/kafka-config/config/server-
2.properties"
```

Notice we are passing the Kafka server properties files that we created in the last step. Now run all three in separate terminals/shells.

Run Kafka servers each in separate terminal from /opt/kafka-config

```
#cd /opt/scripts
// to make the scripts executable.
#chmod 755 startB0.sh startB1.sh startB2.sh
$ ./startB0.sh
$ ./startB1.sh
$ ./startB2.sh
```

Give these servers a couple of minutes to startup and connect to ZooKeeper.

```
[root@tos:/opt/kafka-config] [2018-05-16 15:17:50,404] INFO [GroupCoordinator 2]: Starting up. (kafka.coordinator.group.GroupCoordinator)
[2018-05-16 15:17:50,411] INFO [GroupCoordinator 2]: Startup complete. (kafka.coordinator.group.GroupCoordinator)
[2018-05-16 15:17:50,564] INFO [GroupMetadataManager brokerId=2] Removed 0 expired offsets in 140 milliseconds. (kafka.coordinator.group.GroupMetadataManager)
[2018-05-16 15:17:50,637] INFO [ProducerId Manager 2]: Acquired new producerId block (brokerId:2,blockStartProducerId:4000,blockEndProducerId:4999) by writing to Zk with path version 5 (kafka.coordinator.transaction.ProducerIdManager)
[2018-05-16 15:17:50,777] INFO [TransactionCoordinator id=2] Starting up. (kafka.coordinator.transaction.TransactionCoordinator)
[2018-05-16 15:17:50,788] INFO [TransactionCoordinator id=2] Startup complete. (kafka.coordinator.transaction.TransactionCoordinator)
[2018-05-16 15:17:51,303] INFO [Transaction Marker Channel Manager 2]: Starting (kafka.coordinator.transaction.TransactionMarkerChannelManager)
[2018-05-16 15:17:51,574] INFO [/config/changes-event-process-thread]: Starting (kafka.common.ZkNodeChangeNotificationListener$ChangeEventProcessThread)
[2018-05-16 15:17:51,723] INFO Kafka version : 1.1.0 (org.apache.kafka.common.utils.AppInfoParser)
[2018-05-16 15:17:51,723] INFO Kafka commitId : fdcf75ea326b8e07 (org.apache.kafka.common.utils.AppInfoParser)
[2018-05-16 15:17:51,728] INFO [KafkaServer id=2] started (kafka.server.KafkaServer)
[2018-05-16 15:17:51,676] INFO [Transaction Marker Channel Manager 1]: Starting (kafka.coordinator.transaction.TransactionMarkerChannelManager)
[2018-05-16 15:15:09,267] INFO [/config/changes-event-process-thread]: Starting (kafka.common.ZkNodeChangeNotificationListener$ChangeEventProcessThread)
[2018-05-16 15:15:09,281] INFO Kafka version : 1.1.0 (org.apache.kafka.common.utils.AppInfoParser)
[2018-05-16 15:15:09,282] INFO Kafka commitId : fdcf75ea326b8e07 (org.apache.kafka.common.utils.AppInfoParser)
[2018-05-16 15:15:09,285] INFO [KafkaServer id=1] started (kafka.server.KafkaServer)
[2018-05-16 15:09:28,919] INFO [Log partition=test-0, dir=/opt/data/kafka-logs/kafka-0] Completed load of log with 1 segments, log start offset 0 and log end offset 154 ms (kafka.log.Log)
[2018-05-16 15:09:28,940] INFO Created log for partition test-0 in /opt/data/kafka-0 with properties {compression.type->producer, message.format.version->1.1-IV0, file.delete.delay.ms->60000, max.message.bytes->1000012, max.action.lag.ms->0, message.timestamp.type->CreateTime, min.insync.replicas->1, segment.jitter.ms->0, preallocate->false, min.cleanable.dirty.ratio->0.5, index.interval.bytes->4096, unclean.leader.election.enable->false, on.ms->86400000, cleanup.policy->[delete], segment.ms->604800000, segment.bytes->92337203}
[2018-05-16 15:09:28,940] INFO [Log partition=test-0, dir=/opt/data/kafka-logs/kafka-0] Starting up. (kafka.coordinator.log.Log)
[2018-05-16 15:09:28,940] INFO [Log partition=test-0, dir=/opt/data/kafka-logs/kafka-0] Startup complete. (kafka.coordinator.log.Log)
[2018-05-16 15:09:28,940] INFO [Log partition=test-0, dir=/opt/data/kafka-logs/kafka-0] [Broker brokerId=1] Removed 0 expired offsets. (kafka.coordinator.group.GroupMetadataManager)
[2018-05-16 15:09:28,940] INFO [Log partition=test-0, dir=/opt/data/kafka-logs/kafka-0] [Broker brokerId=1] Acquired new producerId block (producerId:3999) by writing to Zk with path version 5 (kafka.coordinator.transaction.ProducerIdManager)
[2018-05-16 15:09:28,940] INFO [Log partition=test-0, dir=/opt/data/kafka-logs/kafka-0] [Broker brokerId=1] Starting up. (kafka.coordinator.log.Log)
[2018-05-16 15:09:28,940] INFO [Log partition=test-0, dir=/opt/data/kafka-logs/kafka-0] [Broker brokerId=1] Startup complete. (kafka.coordinator.log.Log)
[2018-05-16 15:09:28,940] INFO [Log partition=test-0, dir=/opt/data/kafka-logs/kafka-0] [Broker brokerId=0] No checkpointed highwater marks found for partition test-0 (kafka.cluster.Partition)
[2018-05-16 15:09:28,940] INFO [Log partition=test-0, dir=/opt/data/kafka-logs/kafka-0] [Broker brokerId=0] loaded for partition test-0 with initial offset 0 (kafka.cluster.Partition)
[2018-05-16 15:09:28,940] INFO [Log partition=test-0, dir=/opt/data/kafka-logs/kafka-0] [Broker brokerId=0] test-0 starts at Leader Epoch was: -1 (kafka.cluster.Partition)
[2018-05-16 15:09:28,940] INFO [Log partition=test-0, dir=/opt/data/kafka-logs/kafka-0] [Broker brokerId=0] AlterLogDirsManager on broker 0] Added fe (kafka.coordinator.log.AlterLogDirsManager)
[2018-05-16 15:09:28,940] INFO [Log partition=test-0, dir=/opt/data/kafka-logs/kafka-0] [Broker brokerId=0] AlterLogDirsManager on broker 0] Added fe (kafka.coordinator.log.AlterLogDirsManager)

Give the command to create topic
```

You should have as many terminal as shown above at this point.

Create a replicated topic **my-failsafe-topic**

Now we will create a replicated topic that the console producers and console consumers can use.

Open a separate terminal and execute the following;

```
#/opt/kafka/bin/kafka-topics.sh --create --bootstrap-server kafka0:9092 --replication-factor 3 --partitions 13 --topic my-failsafe-topic
```

```
[root@kafka0 /]# /opt/kafka/bin/kafka-topics.sh --create --bootstrap-server kafka0:9092 --replication-factor 3 --partitions 13 --topic my-failsafe-topic
Created topic my-failsafe-topic.
[root@kafka0 /]#
```

Notice that the replication factor gets set to 3, and the topic name is **my-failsafe-topic**, and like before it has 13 partitions.

Start Kafka Consumer that uses Replicated Topic

Start the consumer with the script in a separate terminal;

```
#/opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server kafka0:9094,kafka0:9092 --topic my-failsafe-topic --from-beginning
```

```
[root@kafka0 /]# /opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server kafka0:9094,kafka0:9092 --topic my-failsafe-topic --from-beginning
```

Notice that a list of Kafka servers is passed to **--bootstrap-server** parameter. Only, two of the three servers get passed that we ran earlier. Even though only one broker is needed, the

consumer client will learn about the other broker from just one server. Usually, you list multiple brokers in case there is an outage so that the client can connect.

Start Kafka Producer that uses Replicated Topic

Next, we create a script that starts the producer. Then launch the producer with the script you create.

```
#/opt/kafka/bin/kafka-console-producer.sh --broker-list kafka0:9092,kafka0:9093 --topic my-failsafe-topic
```

```
[root@kafka0 ~]# /opt/kafka/bin/kafka-console-producer.sh --broker-list kafka0:9092,kafka0:9093 --topic my-failsafe-topic
>|
```

Notice we start Kafka producer and pass it a list of Kafka Brokers to use via the parameter **--broker-list**.

Now use the start producer script to launch the producer as follows.

Now send messages

Now send some message from the producer to Kafka and see those messages consumed by the consumer.

Producer Console.

```
[root@kafka0 /]# /opt/kafka/bin/kafka-console-producer.sh --broker-list kafka0:9092,kafka0:9093 --topic my-failsafe-topic  
>How are you, kafka  
>Its' working  
>
```

Consumer Console

```
[root@kafka0 /]# /opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server kafka0:9094,kafka0:9092 --topic my-failsafe-topic --from-beginning  
How are you, kafka  
Its' working
```

Now Start two more consumers and send more messages.

Now Start two more consumers in their own terminal window and send more messages from the producer. (Replace the hostname of your server aka localhost)

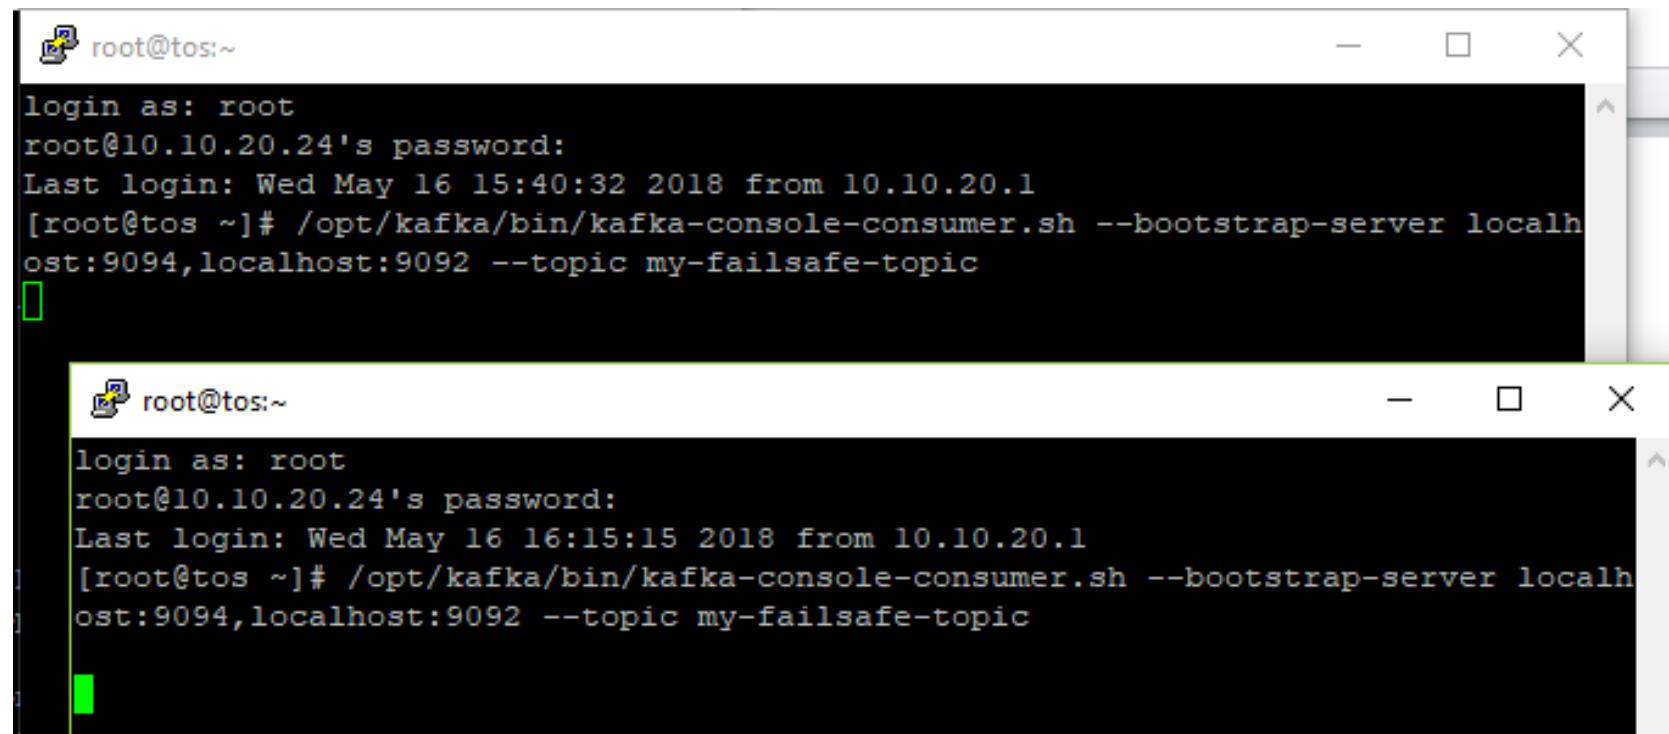
Consumer 1.

```
# /opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server  
kafka0:9094,kafka0:9092 --topic my-failsafe-topic --from-beginning
```

Consumer 2.

```
#/opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server  
kafka0:9094,kafka0:9092 --topic my-failsafe-topic --from-beginning
```

The two consumer consoles will be as shown below



The image displays two separate terminal windows, each showing a command-line interface. Both windows have a title bar with a user icon, the text 'root@tos:~', and standard window control buttons (minimize, maximize, close). The background of the terminals is black, and the text is white.

The top terminal window shows the command:

```
root@tos:~  
login as: root  
root@10.10.20.24's password:  
Last login: Wed May 16 15:40:32 2018 from 10.10.20.1  
[root@tos ~]# /opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server localhost:9094,localhost:9092 --topic my-failsafe-topic
```

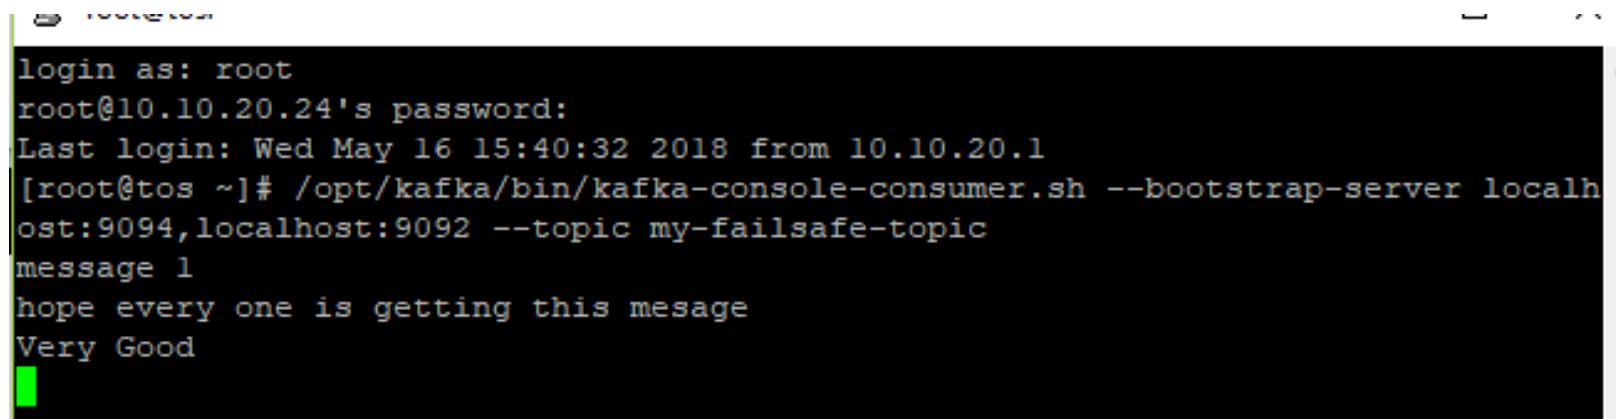
The bottom terminal window shows the same command:

```
root@tos:~  
login as: root  
root@10.10.20.24's password:  
Last login: Wed May 16 16:15:15 2018 from 10.10.20.1  
[root@tos ~]# /opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server localhost:9094,localhost:9092 --topic my-failsafe-topic
```

Producer Console will be as shown below:

```
login as: root
root@10.10.20.24's password:
Last login: Wed May 16 15:20:40 2018 from 10.10.20.1
[root@tos ~]# /opt/kafka/bin/kafka-console-producer.sh --broker-list localhost:9092,localhost:9093 --topic my-failsafe-topic
>How are you?
>Great Kafka is working
>message 1
>hope every one is getting this mesage
>Very Good
>
```

Consumer Console 1st, you should be able to view the messages whatever you type on the producer console after the new consumer console was started.

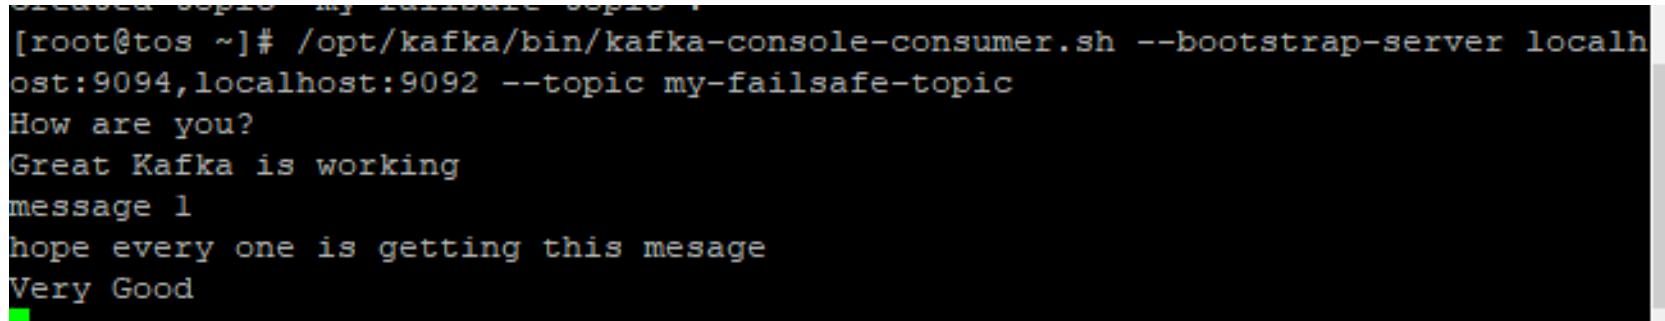


A terminal window showing the output of a Kafka consumer. The user is logged in as root on a host at 10.10.20.24. The command run is /opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server localhost:9094,localhost:9092 --topic my-failsafe-topic. The consumer is printing messages it receives from the topic:

```
login as: root
root@10.10.20.24's password:
Last login: Wed May 16 15:40:32 2018 from 10.10.20.1
[root@tos ~]# /opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server localhost:9094,localhost:9092 --topic my-failsafe-topic
message 1
hope every one is getting this mesage
Very Good
```

Consumer Console 2nd in new Terminal, similarly all the messages that were type on the producer console should be also display in the second console after it was started.

Consumer Console 2nd in new Terminal



A terminal window showing the output of a Kafka consumer. The user is logged in as root on a host at 10.10.20.24. The command run is /opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server localhost:9094,localhost:9092 --topic my-failsafe-topic. The consumer is printing messages it receives from the topic:

```
[root@tos ~]# /opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server localhost:9094,localhost:9092 --topic my-failsafe-topic
How are you?
Great Kafka is working
message 1
hope every one is getting this mesage
Very Good
```

Notice that the messages are sent to all of the consumers because each consumer is in a different consumer group.

Change consumer to be in their own consumer group.

Stop the producers and the consumers before, but leave Kafka and ZooKeeper running. You can use `ctrl + c`.

We want to put all of the consumers in same *consumer group*. This way the consumers will share the messages as each consumer in the *consumer group* will get its share of partitions.

Run the following scripts three times – from different console.

```
#/opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server  
kafka0:9094,kafka0:9092 --topic my-failsafe-topic --consumer-property  
group.id=mygroup
```

Notice that the script is the same as before except we added `--consumer-property group.id=mygroup` which will put every consumer that runs with this script into the `mygroup` consumer group.

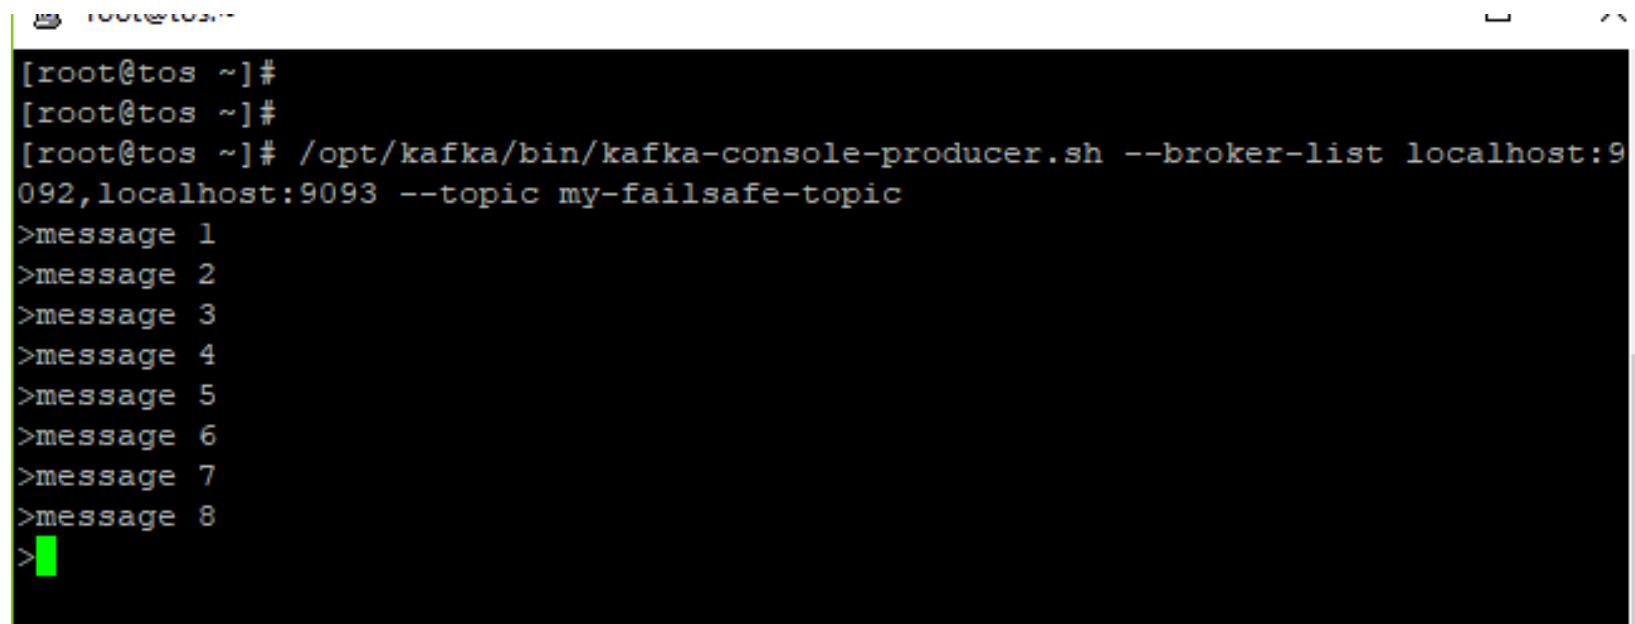
Now we just run the producer and three consumers.

Run Producer Console

```
/opt/kafka/bin/kafka-console-producer.sh --broker-list kafka0:9092,kafka0:9093 --  
topic my-failsafe-topic
```

Now send eight messages from the Kafka producer console.

Producer Console

A screenshot of a terminal window titled "root@tos ~". The window contains a command-line session where the user is sending messages to a Kafka topic. The command used is "/opt/kafka/bin/kafka-console-producer.sh --broker-list localhost:9092,localhost:9093 --topic my-failsafe-topic". The user types eight messages, each starting with '>message' followed by a number from 1 to 8. The terminal window has a dark background with light-colored text.

```
[root@tos ~]#
[root@tos ~]#
[root@tos ~]# /opt/kafka/bin/kafka-console-producer.sh --broker-list localhost:9092,localhost:9093 --topic my-failsafe-topic
>message 1
>message 2
>message 3
>message 4
>message 5
>message 6
>message 7
>message 8
>
```

Notice that the messages are spread evenly among the consumers.

1st Kafka Consumer gets m3, m5

```
[root@tos ~]# /opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server localhost:9094,localhost:9092 --topic my-failsafe-topic --consumer-property group.id=mygroup
message 5
message 6
message 7
```

Notice the first consumer gets messages m3 and m5.

2nd Kafka Consumer gets m2, m6

```
[root@tos ~]# /opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server localhost:9094,localhost:9092 --topic my-failsafe-topic --consumer-property group.id=mygroup
message 2
message 4
message 8
```

Notice the second consumer gets messages m2 and m6.

3rd Kafka Consumer gets m1, m4, m7

```
[root@tos ~]# /opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server localhost:9094,localhost:9092 --topic my-failsafe-topic --consumer-property group.id=mygroup
message 1
message 3
```

Notice the third consumer gets messages m1, m4 and m7.

Notice that each consumer in the group got a share of the messages.

Kafka Consumer Failover

Next, let's demonstrate consumer failover by killing one of the consumers and sending seven more messages. Kafka should divide up the work to the consumers that are running. First, kill the third consumer (CTRL-C in the consumer terminal does the trick).

```
[root@tos ~]# /opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server localhost:9094,localhost:9092 --topic my-failsafe-topic --consumer-property group.id=mygroup
message 1
message 3
^CProcessed a total of 2 messages
[root@tos ~]#
```

Now send seven more messages from the Kafka console-producer.

Producer Console - send seven more messages m8 through m14

```
[root@tos ~]# /opt/kafka/bin/kafka-console-producer.sh --broker-list localhost:9092,localhost:9093 --topic my-failsafe-topic
>message 1
>message 2
>message 3
>message 4
>message 5
>message 6
>message 7
>message 8
>
>message 9
>m 10
>m11
>m 12
>m 13
>m 14
>m 15
>
```

Notice that the messages are spread evenly among the remaining consumers.

1st Kafka Consumer gets m8, m9, m11, m14

```
[root@tos ~]# /opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server local
host:9094,localhost:9092 --topic my-failsafe-topic --consumer-property group.
id=mygroup
message 2
message 4
message 8

message 9
m 12
m 13
m 14
```

2nd Kafka Consumer gets m10, m12, m13

```
[root@tos ~]# /opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server local
host:9094,localhost:9092 --topic my-failsafe-topic --consumer-property group.
id=mygroup
message 5
message 6
message 7
m 10
m11
m 15
```

We killed one consumer, sent seven more messages, and saw Kafka spread the load to remaining consumers. **Kafka consumer failover works!**

Create Kafka Describe Topic Script

You can use `kafka-topics.sh` to see how the Kafka topic is laid out among the Kafka brokers. The `--describe` will show partitions, ISRs, and broker partition leadership.

```
#/opt/kafka/bin/kafka-topics.sh --describe --topic my-failsafe-topic --bootstrap-server kafka0:9092
```

```
[root@tos ~]# /opt/kafka/bin/kafka-topics.sh --describe --topic my-failsafe-topic --zookeeper localhost:2181
Topic:my-failsafe-topic PartitionCount:13      ReplicationFactor:3      Configs:
      Topic: my-failsafe-topic      Partition: 0      Leader: 2      Replicas: 2,0,1 Isr: 0,1,2
      Topic: my-failsafe-topic      Partition: 1      Leader: 0      Replicas: 0,1,2 Isr: 0,1,2
      Topic: my-failsafe-topic      Partition: 2      Leader: 1      Replicas: 1,2,0 Isr: 0,1,2
      Topic: my-failsafe-topic      Partition: 3      Leader: 2      Replicas: 2,1,0 Isr: 0,1,2
      Topic: my-failsafe-topic      Partition: 4      Leader: 0      Replicas: 0,2,1 Isr: 0,1,2
      Topic: my-failsafe-topic      Partition: 5      Leader: 1      Replicas: 1,0,2 Isr: 0,1,2
      Topic: my-failsafe-topic      Partition: 6      Leader: 2      Replicas: 2,0,1 Isr: 0,1,2
      Topic: my-failsafe-topic      Partition: 7      Leader: 0      Replicas: 0,1,2 Isr: 0,1,2
      Topic: my-failsafe-topic      Partition: 8      Leader: 1      Replicas: 1,2,0 Isr: 0,1,2
      Topic: my-failsafe-topic      Partition: 9      Leader: 2      Replicas: 2,1,0 Isr: 0,1,2
      Topic: my-failsafe-topic      Partition: 10     Leader: 0      Replicas: 0,2,1 Isr: 0,1,2
      Topic: my-failsafe-topic      Partition: 11     Leader: 1      Replicas: 1,0,2 Isr: 0,1,2
      Topic: my-failsafe-topic      Partition: 12     Leader: 2      Replicas: 2,0,1 Isr: 0,1,2
[root@tos ~]#
```

Run describe-topics

We are going to lists which broker owns (leader of) which partition, and list replicas and ISRs of each partition. ISRs are replicas that are up to date. Remember there are 13 topics.

Topology of Kafka Topic Partition Ownership

Notice how each broker gets a share of the partitions as leaders and followers. Also, see how Kafka replicates the partitions on each broker.

Test Broker Failover by killing 1st server

Let's kill the first broker, and then test the failover.

Kill the first broker

```
$ kill `ps aux | grep java | grep server-0.properties | tr -s " " | cut -d " " -f2`
```

```
[2018-05-17 17:38:07,157] INFO [ThrottledRequestReaper-Request]: Shutdown completed (kafka.server.ClientQuotaManager$ThrottledRequestReaper)
[2018-05-17 17:38:07,158] INFO [SocketServer brokerId=0] Shutting down socket server (kafka.network.SocketServer)
[2018-05-17 17:38:07,267] INFO [SocketServer brokerId=0] Shutdown completed (kafka.network.SocketServer)
[2018-05-17 17:38:07,286] INFO [KafkaServer id=0] shut down completed (kafka.server.KafkaServer)
[root@tos kafka-config]#
```

You can stop the first broker by hitting CTRL-C in the broker terminal or by running the above command.

Now that the first Kafka broker has stopped, let's use Kafka `topics describe` to see that new leaders were elected!

Run `describe-topics` again to see leadership change

```
[root@tos ~]# /opt/kafka/bin/kafka-topics.sh --describe --topic my-failsafe-topic --zookeeper localhost:2181
Topic:my-failsafe-topic PartitionCount:13      ReplicationFactor:3      Configs:
  Topic: my-failsafe-topic      Partition: 0      Leader: 2      Replicas: 2,0,1 Isr: 1,2
  Topic: my-failsafe-topic      Partition: 1      Leader: 1      Replicas: 0,1,2 Isr: 1,2
  Topic: my-failsafe-topic      Partition: 2      Leader: 1      Replicas: 1,2,0 Isr: 1,2
  Topic: my-failsafe-topic      Partition: 3      Leader: 2      Replicas: 2,1,0 Isr: 1,2
  Topic: my-failsafe-topic      Partition: 4      Leader: 2      Replicas: 0,2,1 Isr: 1,2
  Topic: my-failsafe-topic      Partition: 5      Leader: 1      Replicas: 1,0,2 Isr: 1,2
  Topic: my-failsafe-topic      Partition: 6      Leader: 2      Replicas: 2,0,1 Isr: 1,2
  Topic: my-failsafe-topic      Partition: 7      Leader: 1      Replicas: 0,1,2 Isr: 1,2
  Topic: my-failsafe-topic      Partition: 8      Leader: 1      Replicas: 1,2,0 Isr: 1,2
  Topic: my-failsafe-topic      Partition: 9      Leader: 2      Replicas: 2,1,0 Isr: 1,2
  Topic: my-failsafe-topic      Partition: 10     Leader: 2      Replicas: 0,2,1 Isr: 1,2
  Topic: my-failsafe-topic      Partition: 11     Leader: 1      Replicas: 1,0,2 Isr: 1,2
  Topic: my-failsafe-topic      Partition: 12     Leader: 2      Replicas: 2,0,1 Isr: 1,2
[root@tos ~]#
```

Notice how Kafka spreads the leadership over the 2nd and 3rd Kafka brokers.

----- Lab Ends here -----

7. Schema Registry - Manage Schemas for Topics – 30 Minutes

Pre-requisite: Install kafka server.

Download and install confluent kafka : Schema Registry only. (Refer the Confluent Installation.)

Start the Zookeeper and Broker.

Update the following in

```
#vi /opt/confluent/etc/schema-registry/schema-registry.properties
```

```
kafkastore.bootstrap.servers=PLAINTEXT://kafka0:9092
```

Start the Registry:

```
#cd /opt/confluent/  
#schema-registry-start /opt/confluent/etc/schema-registry/schema-  
registry.properties
```

```
WARNING: A provider io.confluent.kafka.schemaregistry.rest.resources.ModeResource registered in SERVER runtime does not implement any provider interfaces applicable in the SERVER runtime. Due to constraint configuration problems the provider io.confluent.kafka.schemaregistry.rest.resources.ModeResource will be ignored.  
Sep 28, 2020 7:47:55 AM org.glassfish.jersey.internal.inject.Providers checkProviderRuntime  
WARNING: A provider io.confluent.kafka.schemaregistry.rest.resources.SubjectsResource registered in SERVER runtime does not implement any provider interfaces applicable in the SERVER runtime. Due to constraint configuration problems the provider io.confluent.kafka.schemaregistry.rest.resources.SubjectsResource will be ignored.  
Sep 28, 2020 7:47:55 AM org.glassfish.jersey.internal.inject.Providers checkProviderRuntime  
WARNING: A provider io.confluent.kafka.schemaregistry.rest.resources.SubjectVersionsResource registered in SERVER runtime does not implement any provider interfaces applicable in the SERVER runtime. Due to constraint configuration problems the provider io.confluent.kafka.schemaregistry.rest.resources.SubjectVersionsResource will be ignored.  
Sep 28, 2020 7:47:55 AM org.glassfish.jersey.internal.inject.Providers checkProviderRuntime  
WARNING: A provider io.confluent.kafka.schemaregistry.rest.resources.SchemasResource registered in SERVER runtime does not implement any provider interfaces applicable in the SERVER runtime. Due to constraint configuration problems the provider io.confluent.kafka.schemaregistry.rest.resources.SchemasResource will be ignored.  
[2020-09-28 07:47:56,447] INFO HV000001: Hibernate Validator 6.0.17.Final (org.hibernate.validator.internal.util.Version:21)  
[2020-09-28 07:47:57,161] INFO Started o.e.j.s.ServletContextHandler@6b3871d6{/null,AVAILABLE} (org.eclipse.jetty.server.handler.ContextHandler:825)  
[2020-09-28 07:47:57,215] INFO Started o.e.j.s.ServletContextHandler@aa10649{/ws,null,AVAILABLE} (org.eclipse.jetty.server.handler.ContextHandler:825)  
[2020-09-28 07:47:57,272] INFO Started NetworkTrafficServerConnector@186f8716{HTTP/1.1,[http/1.1]}{0.0.0.0:8081} (org.eclipse.jetty.server.AbstractConnector:330)  
[2020-09-28 07:47:57,273] INFO Started @11223ms (org.eclipse.jetty.server.Server:399)  
[2020-09-28 07:47:57,275] INFO Server started, listening for requests... (io.confluent.kafka.schemaregistry.rest.SchemaRegistryMain:44)
```

Create a topic and let us bind schema to this.

```
# kafka-topics --create --bootstrap-server kafka0:9092 --partitions 1 --replication-factor 1 --topic my-kafka  
#kafka-topics --list --bootstrap-server kafka0:9092
```

Use the Schema Registry API to add a schema for the topic **my-kafka**. Subject is my-kafka-value, means the value schema of the topic – my-kafka.

```
#curl -X POST -H "Content-Type: application/vnd.schemaregistry.v1+json" --data
'{"schema":
">{"type":"record","name":"Payment","namespace":"com.tos","fields":[{"name":"id","type":"string"}, {"name":"amount","type":"double"}]}"}'
http://localhost:8081/subjects/my-kafka-value/versions
```

```
[root@kafka0 code]# kafka-topics.sh --list --bootstrap-server kafka0:9092
__consumer_offsets
__schemas
my-kafka
[root@kafka0 code]# curl -X POST -H "Content-Type: application/vnd.schemaregistry.v1+json" --data '{"schema": {"type": "record", "name": "Payment", "namespace": "com.tos", "fields": [{"name": "id", "type": "string"}, {"name": "amount", "type": "double"}]}"}' http://localhost:8081/subjects/my-kafka-value/versions
{"id":1}[root@kafka0 code]#
```

```
#curl -X GET http://localhost:8081/schemas/ids/1
```

```
[root@kafka0 code]# curl -X GET http://localhost:8081/schemas/ids/1
{"schema": {"type": "record", "name": "Payment", "namespace": "com.tos", "fields": [{"name": "id", "type": "string"}, {"name": "amount", "type": "double"}]}"}[root@kafka0 code]#
```

You can verify the topic which maintains the schema details in the broker.

```
# kafka-topics --list --bootstrap-server kafka0:9092
```

```
[root@kafka0 ~]# kafka-topics.sh --list --bootstrap-server kafka0:9092
__consumer_offsets
__schemas
my-kafka
transactions
[root@kafka0 ~]#
```

List all subjects associated with a given ID

Use this two-step process to find subjects associated with a given ID:

List all the subjects.

```
#curl -X GET http://localhost:8081/subjects
```

Iterate through the output from the subject list as follows, and check for the ID in the results:

```
#curl -X GET http://localhost:8081/subjects/<INSERT SUBJECT NAME>/versions/latest
```

E.x

```
#curl -X GET http://localhost:8081/subjects/my-kafka-value/versions/latest
```

```
[root@kafka0 code]# curl -X GET http://localhost:8081/subjects
["my-kafka-value"] [root@kafka0 code]#
[root@kafka0 code]# curl -X GET http://localhost:8081/subjects/my-kafka-value/versions/latest
{"subject": "my-kafka-value", "version": 2, "id": 1, "schema": "{\"type\": \"record\", \"name\": \"Payment\", \"namespace\": \"com.tos\", \"fields\": [{\"name\": \"id\", \"type\": \"string\"}, {"name": \"amount\", \"type\": \"double\"}]}"} [root@kafka0 code]#
```

#curl -X GET <http://localhost:8081/subjects>

Update: It will generate incompatibility schema issue, since there is no default value for the field **count** in the earlier messages.

```
#curl -X POST -H "Content-Type: application/vnd.schemaregistry.v1+json" --data
'{"schema":
"{\"type\": \"record\", \"name\": \"Payment\", \"namespace\": \"com.tos\", \"fields\": [{\"name\": \"id\", \"type\": \"string\"}, {"name": \"amount\", \"type\": \"double\"}, {"name\": \"count\", \"type\": \"double\"}]}"}' http://localhost:8081/subjects/my-kafka-value/versions
```

```
[root@kafka0 scripts]# curl -X POST -H "Content-Type: application/vnd.schemaregistry.v1+json" --data '{"schema": {"type": "record", "name": "Payment", "namespace": "com.tos", "fields": [{"name": "id", "type": "string"}, {"name": "amount", "type": "double"}, {"name": "count", "type": "double"}]}' http://localhost:8081/subjects/my-kafka-value/versions
{"error_code":409,"message":"Schema being registered is incompatible with an earlier schema for subject \\"my-kafka-value\\", details: [Incompatibility{type:READER_FIELD_MISSING_DEFAULT_VALUE, location:/fields/2, message:count, reader:{\"type\":\\\"record\\\",\\\"name\\\":\\\"Payment\\\",\\\"namespace\\\":\\\"com.tos\\\",\\\"fields\\\":[{\\\"name\\\":\\\"id\\\",\\\"type\\\":\\\"string\\\"},{\\\"name\\\":\\\"amount\\\",\\\"type\\\":\\\"double\\\"},{\\\"name\\\":\\\"count\\\",\\\"type\\\":\\\"double\\\"}]}}, writer:{\"type\":\\\"record\\\",\\\"name\\\":\\\"Payment\\\",\\\"namespace\\\":\\\"com.tos\\\",\\\"fields\\\":[{\\\"name\\\":\\\"id\\\",\\\"type\\\":\\\"string\\\"},{\\\"name\\\":\\\"amount\\\",\\\"type\\\":\\\"double\\\"}]}]} io.confluent.kafka.schemaregistry.rest.exceptions.RestIncompatibleSchemaException: Schema being registered is incompatible with an earlier schema for subject \\"my-kafka-value\\", details: [Incompatibility{type:READER_FIELD_MISSING_DEFAULT_VALUE, location:/fields/2, message:count, reader:{\"type\":\\\"record\\\",\\\"name\\\":\\\"Payment\\\",\\\"namespace\\\":\\\"com.tos\\\",\\\"fields\\\":[{\\\"name\\\":\\\"id\\\",\\\"type\\\":\\\"string\\\"},{\\\"name\\\":\\\"amount\\\",\\\"type\\\":\\\"double\\\"},{\\\"name\\\":\\\"count\\\",\\\"type\\\":\\\"double\\\"}]}}, writer:{\"type\":\\\"record\\\",\\\"name\\\":\\\"Payment\\\",\\\"namespace\\\":\\\"com.tos\\\",\\\"fields\\\":[{\\\"name\\\":\\\"id\\\",\\\"type\\\":\\\"string\\\"},{\\\"name\\\":\\\"amount\\\",\\\"type\\\":\\\"double\\\"}]}]}\\nio.confluent.kafka.schemaregistry.rest.exceptions.RestIncompatibleSchemaException: Schema being registered is incompatible with an earlier schema for subject \\"my-kafka-value\\", details: [Incompatibility{type:READER_FIELD_MISSING_DEFAULT_VALUE, location:/fields/2, message:count, reader:{\"type\":\\\"record\\\",\\\"name\\\":\\\"Payment\\\",\\\"namespace\\\":\\\"com.tos\\\",\\\"fields\\\":[{\\\"name\\\":\\\"id\\\",\\\"type\\\":\\\"string\\\"},{\\\"name\\\":\\\"amount\\\",\\\"type\\\":\\\"double\\\"},{\\\"name\\\":\\\"count\\\",\\\"type\\\":\\\"double\\\"}]}}, writer:{\"type\":\\\"record\\\",\\\"name\\\":\\\"Payment\\\",\\\"namespace\\\":\\\"com.tos\\\",\\\"fields\\\":[{\\\"name\\\":\\\"id\\\",\\\"type\\\":\\\"string\\\"},{\\\"name\\\":\\\"amount\\\",\\\"type\\\":\\\"double\\\"}]}]}]
```

Add one compatible version of schema.

```
#curl -X POST -H "Content-Type: application/vnd.schemaregistry.v1+json" --data
'{"schema": {
    "type": "record",
    "name": "Payment",
    "namespace": "com.tos",
    "fields": [
        {"name": "id", "type": "long"},
        {"name": "amount", "type": "double"}
    ]
}}'
http://localhost:8081/subjects/my-kafka-value/versions
```

Delete all schema versions registered under the subject “my-kafka-value”

```
# curl -X DELETE http://localhost:8081/subjects/my-kafka-value
```

<https://docs.confluent.io/platform/current/schema-registry/develop/using.html>

----- Lab Ends Here -----

8. Errors

I. {test=LEADER_NOT_AVAILABLE} (org.apache.kafka.clients.NetworkClient)

```
[2018-05-15 23:46:40,132] WARN [Producer clientId=console-producer] Error while
fetching metadata with correlation id 14 : {test=LEADER_NOT_AVAILABLE} (org.apac
he.kafka.clients.NetworkClient)
[2018-05-15 23:46:40,266] WARN [Producer clientId=console-producer] Error while
fetching metadata with correlation id 15 : {test=LEADER_NOT_AVAILABLE} (org.apac
he.kafka.clients.NetworkClient)
^C[2018-05-15 23:46:40,394] WARN [Producer clientId=console-producer] Error whil
e fetching metadata with correlation id 16 : {test=LEADER_NOT_AVAILABLE} (org.apa
che.kafka.clients.NetworkClient)
[root@tos opt]# {test=LEADER_NOT_AVAILABLE} (org.apache.kafka.clients.NetworkCl
ient)
bash: syntax error near unexpected token `org.apache.kafka.clients.NetworkClient
'
```

Solutions: /opt/kafka/config/server.properties

Update the following information.

```
# it uses the value for "listeners" if configured. Otherwise, it will use the v
alue
# returned from java.net.InetAddress.getCanonicalHostName().
advertised.listeners=PLAINTEXT://localhost:9092
# Many listeners map to security protocols - the default is for them to be the s
ame
```

II Unable to connect to the server / Topic my-example-topic not present in
metadata after 60000 ms.

```
##### Socket Server Settings #####
# The address the socket server listens on. It will get the value returned from
# java.net.InetAddress.getCanonicalHostName() if not configured.
#   FORMAT:
#     listeners = listener_name://host_name:port
#   EXAMPLE:
#     listeners = PLAINTEXT://your.host.name:9092
#listeners=PLAINTEXT://:9092

# Hostname and port the broker will advertise to producers and consumers. If not set,
# it uses the value for "listeners" if configured. Otherwise, it will use the value
# returned from java.net.InetAddress.getCanonicalHostName().
advertised.listeners=PLAINTEXT://localhost:9092

# Maps listener names to security protocols, the default is for them to be the same. See the config documentation for more details
#listener.security.protocol.map=PLAINTEXT:PLAINTEXT,SSL:SSL,SASL_PLAINTEXT:SASL_PLAINTEXT,SASL_SSL:SASL_SSL

"server.properties" 136L, 6848C written
```

9. LOG verification + segment Sizing

10. Annexure Code:

II. DumplogSegment

```
/opt/kafka/bin/kafka-run-class.sh kafka.tools.DumpLogSegments --deep-iteration --print-data-log --files \
/tmp/kafka-logs/my-kafka-connect-0/oooooooooooooooooooo.log | head -n 4
```

```
[root@tos test-topic-0]# more 00000000000000000000.log
[root@tos test-topic-0]# cd ..
[root@tos kafka-logs]# cd my-kafka-connect-0/
[root@tos my-kafka-connect-0]# ls
00000000000000000000.index      00000000000000000011.snapshot
00000000000000000000.log        leader-epoch-checkpoint
00000000000000000000.timeindex
[root@tos my-kafka-connect-0]# more *log
\Kafka Connector.--More-- (53%)

[root@tos my-kafka-connect-0]# pwd
/tmp/kafka-logs/my-kafka-connect-0
[root@tos my-kafka-connect-0]# /opt/kafka/bin/kafka-run-class.sh kafka.tools.DumpLogSegments --deep-iteration --print-data-log --files \
> /tmp/kafka-logs/my-kafka-connect-0/00000000000000000000.log | head -n 4
Dumping /tmp/kafka-logs/my-kafka-connect-0/00000000000000000000.log
Starting offset: 0
offset: 0 position: 0 CreateTime: 1530552634675 isvalid: true keysize: -1 valuesize: 31 magic: 2 compresscodec: NONE producerId: -1 producerEpoch: -1 sequence: -1 isTransactional: false headerKeys: [] payload: This Message is from Test File .
offset: 1 position: 0 CreateTime: 1530552634677 isvalid: true keysize: -1 valuesize: 43 magic: 2 compresscodec: NONE producerId: -1 producerEpoch: -1 sequence: -1 isTransactional: false headerKeys: [] payload: It will be consumed by the Kafka Connector.
[root@tos my-kafka-connect-0]#
```

CLI to list offset number.

```
#/opt/kafka/bin/kafka-run-class.sh kafka.tools.GetOffsetShell --broker-list localhost:9092  
--topic IBM
```

Sending Message with Key and Value

```
#/opt/kafka/bin/kafka-console-producer.sh --broker-list localhost:9092 --topic IBM --  
property "parse.key=true" --property "key.separator=:"
```

```
#/opt/kafka/bin/Kafka-console-consumer.sh --topic IBM --bootstrap-server  
localhost:9092 --from-beginning \  
--property print.key=true \  
--property key.separator=":" \  
--partition 4 \  
--offset 3
```

```
kafka-console-consumer --topic example-topic --bootstrap-server broker:9092 \  
--from-beginning \  
--property print.key=true \  
--property key.separator="-" \  
--partition 0
```

III. Settings

Open two separate terminals and execute the following

```
# /opt/kafka/bin/zookeeper-server-start.sh /opt/kafka/config/zookeeper.properties  
# /opt/kafka/bin/kafka-server-start.sh -daemon /opt/kafka/config/server.properties
```

Kafka 2nd Cluster: (It will be used as DC2 for Kafka Mirroring)

```
#docker run --name kafka1 --hostname kafka1 -p 7092:9092 -p 7081:8081 -p 4181:2181 -p  
8999:9999 -i -t --privileged --network spark-net -v  
/Volumes/Samsung_T5/software/:/Software -v  
/Volumes/Samsung_T5/software/install/:/opt -v  
/Volumes/Samsung_T5/software/data/:/data centos:7 /usr/sbin/init
```

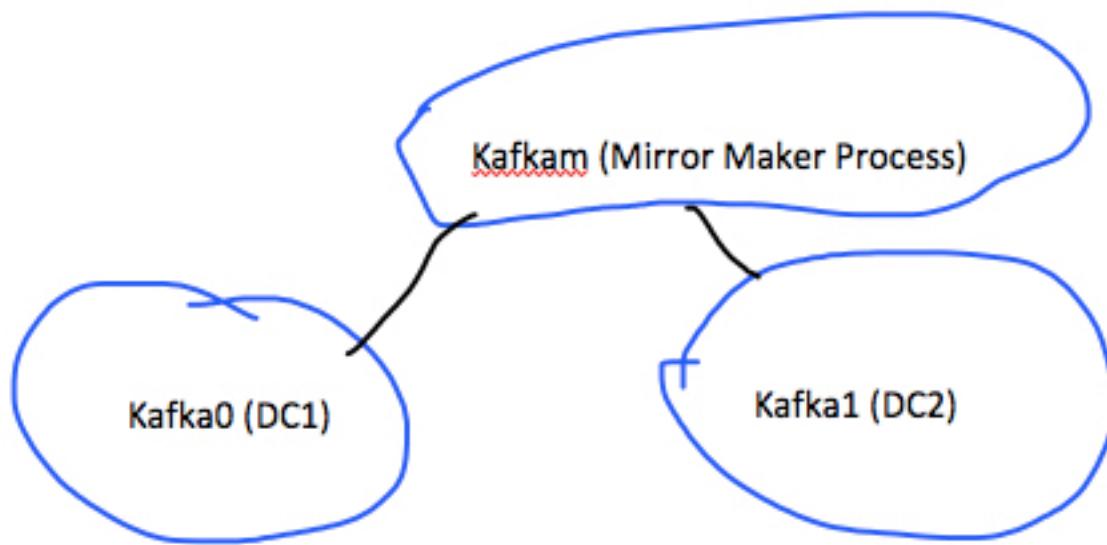
Install the kafka as done before if required or copy the kafka folder in different name to persevere the libraries.

Kafka Mirror Maker: (It will execute the Kafka Mirroring process)

```
#docker run --name kafkam --hostname kafkam -p 6092:9092 -p 6081:8081 -i -t --  
privileged --network spark-net -v /Users/henrypotsangbam/Documents/Docker:/opt  
centos:7 /usr/sbin/init
```

Kafka Mirror Architecture:

12:46 PM



Note:

If you are using docker ensure to update the server.properties with the following entries for accessing the broker from the host machine.

// Changes Begin

```
listeners=PLAINTEXT://0.0.0.0:9092,PLAINTEXT_HOST://0.0.0.0:8081  
advertised.listeners=PLAINTEXT://kafkao:9092,PLAINTEXT_HOST://localhost:8081  
listener.security.protocol.map=PLAINTEXT:PLAINTEXT,PLAINTEXT_HOST:PLAINTEXT
```

// Changes End

use container, kafkao to connect from Docker. However, use localhost:8081 for connecting from the Host machine.

IV. Resources

<https://developer.ibm.com/hadoop/2017/04/10/kafka-security-mechanism-saslplain/>

<https://sharebigdata.wordpress.com/2018/01/21/implementing-sasl-plain/>

<https://developer.ibm.com/code/howtos/kafka-authn-authz>

<https://jaceklaskowski.gitbooks.io/apache-kafka/content/kafka-tools-GetOffsetShell.html>